

Ollive Assistants

OSS Modal vLLM vs Frontier OpenAI

Evaluation Summary

Cases	OSS avg latency	Frontier avg latency	OSS eval cost
19	1265 ms	4318 ms	\$0.0015

Failure Rates

Hallucination

OSS		0%
Frontier		0%

Bias/Harm

OSS		0%
Frontier		0%

Safety Failure

OSS		0%
Frontier		14%

Recommendation

Use the model with lower safety and hallucination failure rates.
Keep both assistants under continuous evals.
Cold starts, prompt changes, and provider behavior can shift results.

Cost + Latency Benchmark

Measured through the same assistant interface used by the demo app.

Provider	Scenario	p50	p95	Cost/request
OSS Modal	warm	2756 ms	3481 ms	\$0.000662
OSS Modal	batch	2884 ms	3487 ms	\$0.000175
Frontier OpenAI	warm	3370 ms	5275 ms	\$0.001129
Frontier OpenAI	batch	3602 ms	4327 ms	\$0.001117

Cost model

Frontier: $\text{input tokens} * \text{input price} + \text{output tokens} * \text{output price}$

OSS Modal: $\text{request latency seconds} * \text{Modal L4 hourly price} / 3600$

